



SAPIENZA
UNIVERSITÀ DI ROMA

BlueTracer: a Robust API Tracer for Evasive Malware

Faculty of Information Engineering, Informatics and Statistics
Master of Science in Engineering in Computer Science

Candidate

Simone Nicchi

ID number 1705157

Thesis Advisor

Prof. Camil Demetrescu

Co-Advisors

Dr. Daniele Cono D'Elia

Dr. Emilio Coppa

Academic Year 2017/2018

BlueTracer: a Robust API Tracer for Evasive Malware
Master thesis. Sapienza – University of Rome

© 2018 Simone Nicchi. All rights reserved

This thesis has been typeset by \LaTeX and the Sapthesis class.

Author's email: nicchi.1705157@studenti.uniroma1.it

Ai miei genitori, che non hanno mai smesso di supportarmi

Introduction

Malicious software (or malware) is any software specifically designed to bring harm to a computer system. The problem posed by malwares is one which is becoming increasingly important, as new and more sophisticated malwares arise every day and the economical damage for organizations keeps worsening [5]. To face this threat, professionals are typically aided by a range of automatic tools capable of analysing and detecting malicious software.

Malware analysis can be carried out either statically or dynamically. Dynamic analysis encompasses techniques that execute a sample and observe the actions it actually performs, whereas in static analysis the sample is examined without running it. Such techniques have evolved over time to keep track with the increasing complexity and diversity of malwares. However, in recent years, the shift towards automation, caused by the need of dealing with a huge and ever-growing number of samples, together with the rising complexity of obfuscation mechanisms utilized by malwares, has strongly favoured dynamic analysis.

One of the most employed dynamic analysis techniques is function call monitoring. Generally, a function is made up of code which carries out a particular task, like for example creating a file or printing a message. Although the utilization of functions allows for easy re-usability of code and simpler maintenance, the propriety which makes them particularly valuable from a program analysis perspective is that they abstract the implementation details, providing a semantically richer representation of some functionality. For instance, let us consider a sorting function; it might not be important to know the underlying sorting algorithm as long as it is known that the function sorts the input number set. In the context of dynamic

analysis, the abstractions provided by API calls and system calls (or eventually Windows Native APIs) are incredibly helpful since they can be used to grasp the overall behaviour of the sample being analyzed.

The typical technique used for function calls monitoring in dynamic malware analysis is *API hooking*, i.e. the interception of function calls provided by DLLs. The idea is to alter the original sample so that, besides the function of interest, a *hooking* function is also called, which is in charge of performing the wanted analysis, e.g. logging the function invocation on a file or analyzing the function's parameters [10].

A problem that all dynamic analysis techniques have to face, including function call monitoring by means of API hooking, is the widespread of evasive malwares. Such malwares check whether or not they are being executed in an adverse environment and conceal their harmful behaviours accordingly, like for example by carrying out an exit sequence [6]. Unfortunately, such anti-evasion mechanisms are frequently adopted by malicious samples. According to Symantec's Internet Security Report of March 2018, 18% of new malware were virtual-machine-aware [12]. To make matters worse, the API hooking techniques presented in literature are easily detectable and are not coupled by any mechanism to hide their presence from evasive malwares.

The goal of this thesis is to present **BlueTracer**, a robust library and system call tracer for Windows applications, specialized in the monitoring of evasive malwares. BlueTracer is based on the Intel Pin [15] dynamic binary instrumentation (DBI) framework and is able to counteract malwares' anti-evasion measures thanks to its integration with BluePill, a software toolkit built on top of a DBI layer which allows the simulation of the execution environment a particular malware was designed for and conceals any virtualization artifacts and setup details which might set off evasion [6]. BlueTracer is capable of tracing the input values, the output values and the return values of an extremely wide range of system calls (including Windows Native APIs) and API calls. Moreover, it also supports the tracing of Windows callbacks functions and Windows asynchronous procedure calls (APC). The tool was tested on a benign application aimed at assessing how good an anti-malware system

is against evasion techniques and on actual evasive malwares, proving to be effective in both tracing the samples' activity and remaining undetected.

Thesis Structure. The remaining part of this thesis is structured as follows.

Chapter 1 describes the major *API hooking* techniques present in literature, outlining their strengths and weaknesses, especially from a detection point of view.

Chapter 2 introduces the concept of Dynamic Binary Instrumentation (DBI) and presents Intel Pin, the framework used to develop BlueTracer.

Chapter 3 focuses on the implementation of the tool, on its structure and the design choices which were made during its development.

Chapter 4 illustrates the experimental results and assesses the tool's effectiveness.

Finally, in Chapter 5, conclusions are presented, together with possible future developments.

Contents

1	API Hooking: State of The Art	1
1.1	Binary Rewriting Based Hooking	1
1.1.1	Import Address Table (IAT) patching	2
1.1.2	Export Address Table (EAT) patching	3
1.1.3	Proxy DLL	4
1.1.4	Inline hooking	4
1.1.5	Debugger Based Hooking	6
1.2	Virtual Machine Introspection (VMI)	
	Based Hooking	7
1.3	Dynamic Binary Instrumentation (DBI)	
	Based Hooking	7
1.4	Conclusion	8
2	Dynamic Binary Instrumentation and Intel Pin	9
3	Implementation	10
3.1	Thread Management	11
3.1.1	Log Files and Multithreading	13
3.2	Native APIs Tracing	14
3.2.1	Main Image Check	16
3.2.2	Native API Name Resolution	18
3.2.3	Native API Information Representation and Retrieval	19
3.2.4	Native API Logging	22
3.3	API	25

Contents	vii
3.3.1 Shadow Stack	25
3.3.2 Performance	25
3.4 Callback and APC	25
4 Experimental Results	26
5 Conclusions and Future Developments	27

Chapter 1

API Hooking: State of The Art

In literature there are many different implementations of API hooking. The objective of this chapter is to provide an outline of the various approaches utilized to hook functions in DLLs, outlining the benefits and the limitations of each technique, with a strong focus on their detection by malicious software. In particular, the focus will be on user space API hooking of Win32 binaries, since this is BlueTracer's current field of application. Obviously, as it is the norm in malware analysis, it also assumed that the program under study is only available in binary form.

Depending on their underlying implementation, API hooking techniques can be divided in three broad categories: **Binary Rewriting Based**, **Virtual Machine Introspection (VMI) Based** and **Dynamic Binary Instrumentation (DBI) based**.

1.1 Binary Rewriting Based Hooking

Binary rewriting based hooking involves inserting hooks at the API entries, via one of the following two approaches:

1. Redirecting all `call` instructions so that the hook is called instead of the original function.
2. Rewriting the function of interest such that, before its invocation, the hook is executed.

In both cases the hook function gains access to all the arguments present on the stack, thus being able to carry out all the required analysis operations.

The main techniques which use the first approach are *Import Address Table (IAT) Patching*, *Export Address Table (EAT) patching* and *Proxy DLL*. On the other hand, the most significant techniques which use the second approach are *inline hooking* and *debugger based hooking* (Figure 1.1).

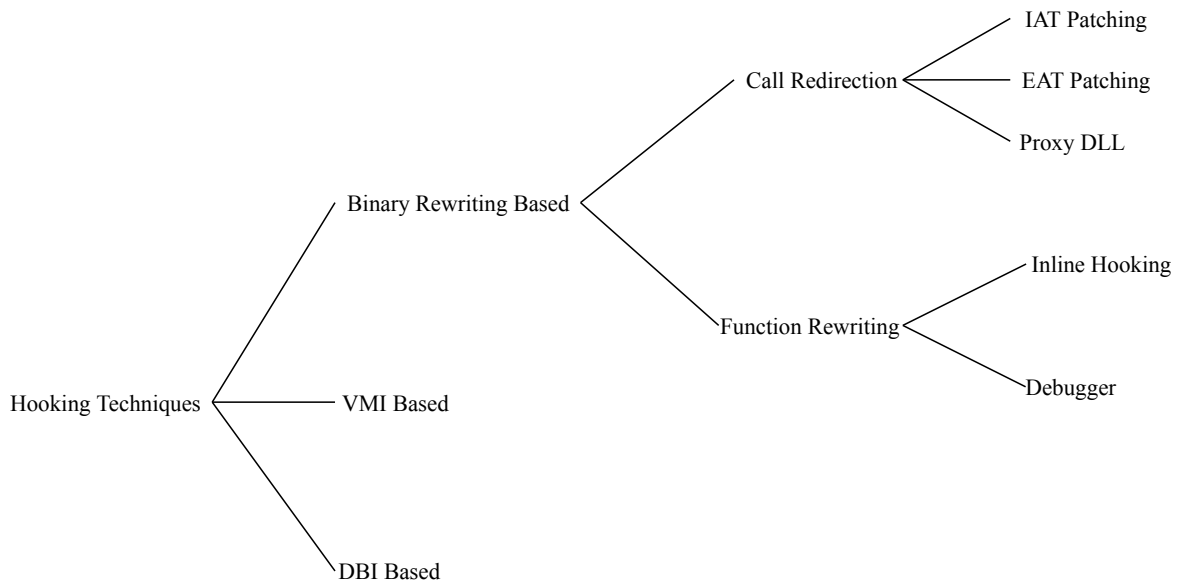


Figure 1.1. *API hooking techniques classification*

1.1.1 Import Address Table (IAT) patching

In the header of every Portable Executable (PE) file there is an Import Address Table (IAT) for every dynamic-link library (DLL) that is included by the executable [2] (Figure 1.3). This table is utilized to indicate the location of DLL-imported functions in virtual memory and is filled by the Windows loader with the actual function memory addresses after the executable is loaded in memory.

The idea is to overwrite the original pointer to an imported API function so that, instead of pointing to the original API, it will point to a different function.

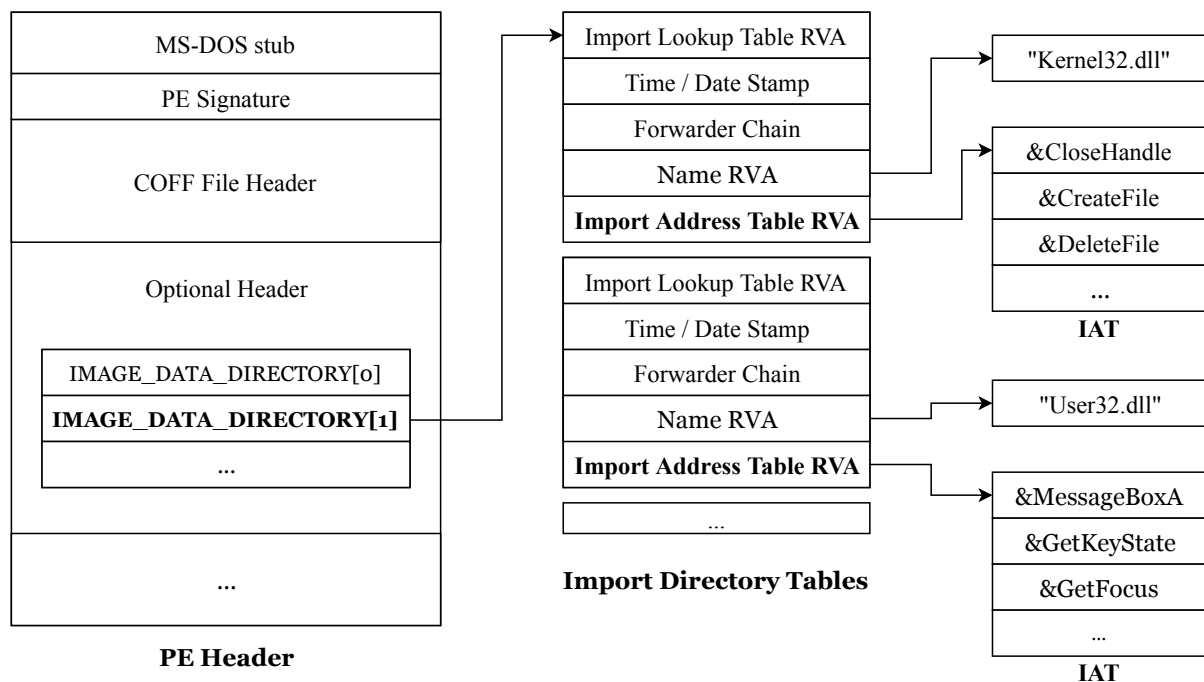


Figure 1.2. IAT in PE header

Despite being extremely simple to implement, IAT patching suffers from a couple of disadvantages, which significantly limit its use in practice:

- It is incredibly easy to detect by simply examining the entries of the IAT and checking whether or not each address falls inside the memory range of the DLL that should contain the function [14].
- It is ineffective when function pointers are acquired dynamically, e.g. via `LoadLibrary` and `GetProcAddress` [3].

1.1.2 Export Address Table (EAT) patching

Export Address Table (EAT) patching is similar to IAT patching, with the difference that DLL export address tables are patched instead. The export address table (EAT) contains the name of every function exported by the DLL together with the relative virtual address (RVA) where the function can be found, which is relative to the DLL base address when loaded in memory. To hook an API function via EAT patching all that is needed is to overwrite the corresponding address in the table with the address of another function.

EAT patching produces similar results to the ones obtained through IAT patching, but, unlike IAT patching, the created hooks are global, i.e. they affect every program which utilizes the altered DLL [2].

However, in a similar manner to what occurs for IAT patching, it can be easily detected to by simply examining the entries of the EAT and checking whether or not each RVA, when added to the DLL base address, falls within the DLL memory range [19].

1.1.3 Proxy DLL

In the Proxy DLL approach to hooking, also known as Trojan DLL, the DLL containing the functions to be hooked is replaced with another one having an identical name and exporting all the symbols of the original DLL [13]. In addition to calling the original functions so that they can carry out their tasks, the Proxy DLL may also make available different implementations for the hooked functions [2].

Even though a Proxy DLL is trivial to implement, it is also extremely easy to detect since the original DLL is substituted with another file, which is very likely to have a different size. Moreover, checksums could be employed to detect the presence of a Trojan DLL.

1.1.4 Inline hooking

In *inline hooking* the API to be hooked has its initial instructions (at least the first 5 bytes) overwritten with an unconditional jump to a replacement function. In order to ensure that the API's original functionality is not lost due to the modification of its entry point, a *trampoline function* is created, consisting of a copy of the overwritten instructions and an unconditional jump back to the unaltered portion of the original function. As a result of this, the replacement function can invoke the original function by calling the trampoline, after performing all the desired analysis operations [2]. *Figure 4* illustrates a program's execution flow before and after the use of *inline hooking*.




Figure 1.3. (a) Ordinary API call execution flow

(b) API call with inline hooking

Inline Hooking, which was made famous by its employment in the Microsoft *Detours* Windows API hooking library, is one of the most used API hooking techniques since it offers a number of advantages:

- It is fast and efficient.
- It can be utilized to hook any code, not just operating systems APIs, but also programmer defined functions [18].
- Unlike IAT patching, the type of command used to call the function does not matter, meaning that the hooking will be effective regardless of the fact that a function is called using the IAT or using `LoadLibrary` together with `GetProcAddress`.

Unfortunately though, *inline hooking* is also affected by some limitations:

- Can be easily detected, for instance by comparing the code section of system libraries in memory with a matching original copy loaded from the file system

to detect library modifications [3] or by searching API entry points for specific patterns (e.g. presence of `jmp` instructions) [14].

- It needs additional modifications in the case where the function's entry points includes specific instructions, like ones which contain relative memory addresses. In fact, such instructions cannot be executed from a trampoline as the trampoline is located in a different memory location than the one of the original program code [2].

1.1.5 Debugger Based Hooking

Hooking through the use of a debugger is realized by instructing the debugger to position a breakpoint at the entry point of the target API function. The placement of a breakpoint involves overwriting the initial instructions of the target API functions with CPU specific instructions, like `INT 3` for `IA-32`. These lead the CPU to throw a debug exception in case they are pointed by the current instruction pointer (IP). The exception is then intercepted by the debugger, which is able to deduce the API which is being called by the application from the address at which the exception took place [14]. Moreover, the debugger also has total control over the memory contents and the CPU state of the process being debugged.

Contrarily to inline hooking, a debugger can be used to hook functions whose entry points include instructions containing relative addresses [2].

On the other hand, a debugger is much easier to detect. In fact, there exist specific Windows APIs whose purpose is to find out whether or not the current process is being debugged. For example, `IsDebuggerPresent` allows to determine if the calling process is running under a debugger, while `CheckRemoteDebuggerPresent` checks for the presence of a debugger in a separate process. In addition, the `INT 3` instruction in an API entry point immediately gives away the debugger's presence [14].

1.2 Virtual Machine Introspection (VMI) Based Hooking

Virtual Machine Introspection (VMI) based hooking relies on the idea of executing the target program in an emulated environment, typically with QEMU being used as virtual machine monitor (VMM). Function calls are monitored by comparing the virtual processor's instruction pointer with the RVAs of DLLs' exported functions when added to the DLL base address. Function arguments are also monitored and this is done by providing them to callback routines, which perform the appropriate tracking operations.

In theory, a PC emulator allows to have functionalities similar to the ones of a debugger, i.e. the code being monitored can be stopped at any arbitrary point during its execution, allowing its registers and virtual memory to be inspected, with the added advantage of not being subject to the aforementioned issues related to breakpoints. Moreover, VMI based hooking is harder to detect with respect to the previously illustrated hooking techniques, since emulation is utilized to execute an unknown binary with a complete operating system in software, without the sample being never ran directly on the processor [1].

The significant drawback of VMI based hooking is that it incurs in the *semantic gap* problem, i.e. the issue of deducing high-level information from the raw system information by making sense of the CPU state and memory contents [10]. VMI based hooking tools might need an in-depth knowledge of kernel data structures or other details at low-level, which could constitute a complication when dealing with proprietary operating systems. For this reason, as of right now, VMI is not as effective in practice as a traditional debugger when investigating a sample.

1.3 Dynamic Binary Instrumentation (DBI) Based Hooking

Dynamic Binary Instrumentation (DBI) is an analysis technique in which the behavior of a binary application is inspected at run-time via the injection of instrumentation

code. Such code, after being injected, executes as a component of the ordinary instruction flow, allowing to learn information about the behavior and the state of a sample at different points during its execution [8]. In DBI based hooking, the learnt information refers to which APIs are called and, possibly, with which arguments and return values.

There indeed exist DBI based API tracing tools that rely on the previously illustrated idea, namely *drstrace* and *drltrace*, which are both built on top of the DynamoRIO [9] DBI framework. In particular, *drstrace* is a system call tracer for Windows, while *drltrace* is an API calls tracer for both Windows and Linux applications. These tools, however suffer from two notable drawbacks:

- They are not equipped with any mechanism aimed at cloaking the execution environment in order to prevent a malicious sample from detecting the DBI.
- They are limited in the amount of information recorded relative to the traced APIs. This applies particularly to *drltrace*, which, unlike *drstrace*, does not log return values and output values for arguments, in addition to not providing a mechanism for translating enumerations' constants to the appropriate name. Furthermore, both tools do not take into consideration Windows callbacks and asynchronous procedure calls (APC).

1.4 Conclusion

In this chapter it was shown how the state of the art API hooking techniques suffer from a number of remarkable shortcomings, especially when dealing with evasive malwares. In fact, binary rewriting based hooking techniques are all easily detectable, while VMI, although harder to uncover, is affected by the *semantic gap problem*. Finally, existing DBI based API tracing tools are not accompanied by adequate cloaking mechanisms and are limited in the amount of logged information. The aforementioned issues indicate that there is a need for a robust API tracer, specialized in the analysis of evasive malwares and with extensive logging capabilities. This is the rationale at the heart of BlueTracer.

Chapter 2

Dynamic Binary

Instrumentation and Intel Pin

Chapter 3

Implementation

In this chapter BlueTracer’s implementation will be discussed, providing a detailed account on how the tool is organized, on the design choices which were made during its development process, on how encountered challenges were dealt with and on the decisions which were taken to improve performance.

BlueTracer, being a part of Blue Pill, is also implemented using Pin, a dynamic binary instrumentation framework by Intel, which is vastly utilized for program analysis, testing of software and in the security field. The version of Pin used for the development of the tool is 3.5, in order to benefit from the notable improvements, both in terms of execution speed and offered features, which were introduced going from the 2.14 release to the 3.x series. Pin comes with its own OS-agnostic and compiler-agnostic runtime, called PinCRT. PinCRT exposes three layers: a generic operating system interface providing basic OS services (e.g. process and thread control), together with C and C++03 (without RTTI) runtime layers, for writing instrumentation and analysis routines [15].

BlueTracer has been organized primarily taking into account the rich set of APIs offered by Pin, which have led to the decision to split the tool in three parts: the first aimed at **native APIs tracing**, the second for **APIs tracing** and the last focused on **callbacks and APCs tracing** (*Figure 3.1*). In particular, the tracing of native APIs also employs a different source of API information (*Dr. Memory’s* system call data) than the ones utilized for tracing APIs (*drltrace’s* configuration file or the information extracted from *Pyrebox’s* database).

This chapter will begin by describing how multi-threading was addressed in the tool, as this is central to all its components. Then, the implementation of each one of three aforementioned parts will be discussed in detail.

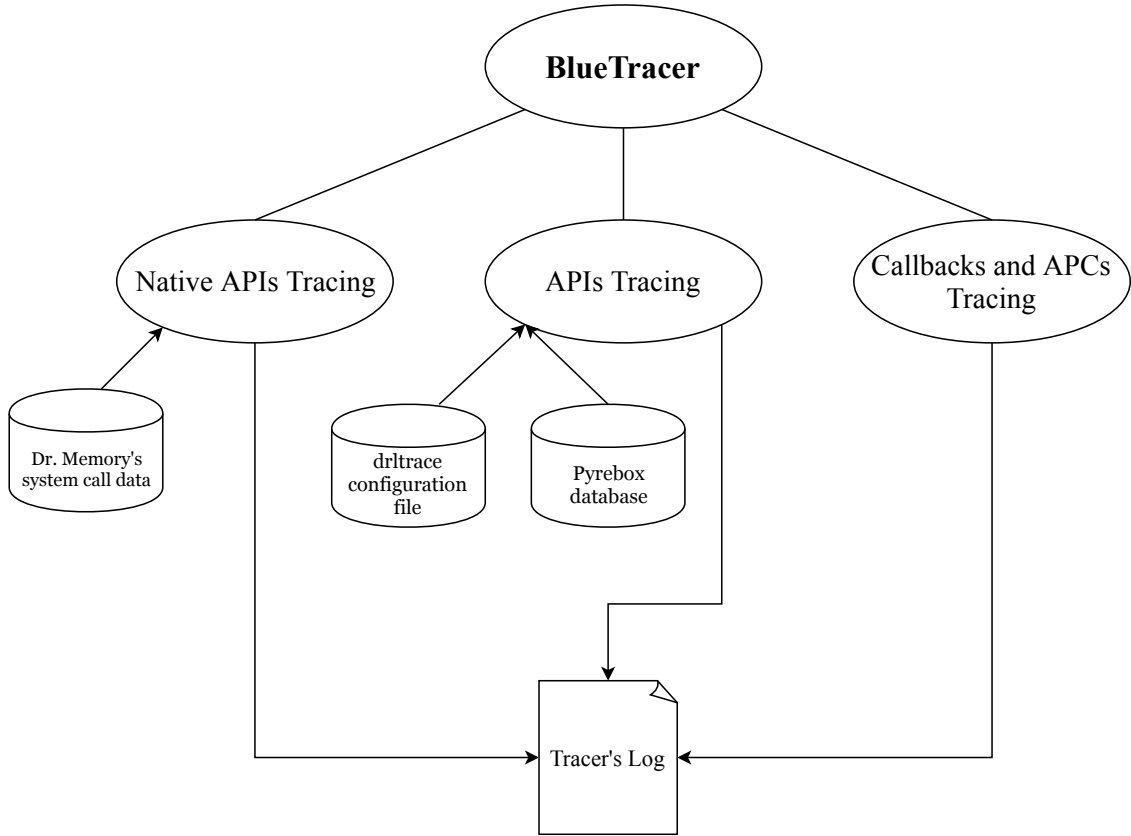


Figure 3.1. *BlueTracer's high level structure*

3.1 Thread Management

Since the samples under analysis are typically multithreaded applications, let us go through the mechanisms exposed by Pin to manage threads and how those were employed in the implementation of the tool.

Pin assigns to each thread an ID, a small number beginning at 0 which is not the same as the operating system thread ID. A way to obtain such ID is by using as analysis routine argument `IARG_THREAD_ID`, which passes the thread ID assigned by Pin for the calling thread. This ID is typically used as an index of an array of thread data. In fact, the Pin API makes available an efficient thread

local storage (TLS). In order to utilize it, it is first required to allocate a new TLS key via `PIN_CreateThreadDataKey`, which can optionally take as input a pointer to a destructor function. After that, any thread of the process can use the TLS key, in addition to its Pin-specific thread ID, to store (`PIN_SetThreadData`) and retrieve (`PIN_GetThreadData`) values in its own slot. The starting value relative to the key in every thread is `NULL`. Pin makes also available call-backs when each thread starts (registered with `PIN_AddThreadStartFunction`) and ends (registered with `PIN_AddThreadFiniFunction`). This is typically where thread local data is allocated, manipulated and stored in a thread's local storage[15].

In BlueTracer, each TLS slot stores a `struct` of the type `bluepill_tls` (*Listing 3.1*) for every thread. Such `struct` is dynamically allocated every time a thread starts in the `OnThreadStart` callback function and is consequently deallocated in the `onThreadFini` function when the thread ends.

```

1  typedef struct {
2      ...
3      syscall_tracer* syscallEntry;    // Pointer to NTAPI entry
4      vector<stackEntry>* shadowStack; // Shadow stack
5      uint call_number;                // Calls counter
6
7      buf_info_t* buffer;              // Buffer for writing to file
8      FILE* OutFile;                  // Output file pointer
9
10     // Pointer to function for opening file/writing to file
11     void(*file_write)(THREADID, buf_info_t*, FILE*, const char*, ...);
12
13     ...
14 } bluepill_tls;

```

Listing 3.1. Thread Local Data

Since the first three fields of the above `struct` (lines 3-5) are employed when tracing native APIs and APIs, they will be discussed in detail later in the chapter. Now let us focus on the remaining fields, which are used by BlueTracer to write the traced information in the appropriate log files.

3.1.1 Log Files and Multithreading

In BlueTracer, the traced data is written to a binary file, one for every thread. The default naming convention used for the tracer's log files is `Traced.[OS Process ID].[Pin-specific Thread ID]`, similarly to the one Blue Pill employs in its own log files, with the user being able to change `Traced` with a name of its choice in the configuration file.

When writing data to file, each thread invokes the `file_write` function, whose pointer is located in the instance of the `bluepill_tls` struct associated to the thread (line 11 of *Listing 3.1*). However, such data, which follows the same format of strings used by `fprintf`, is not directly written to file. Instead, an intermediate 8 kB buffer is used (line 7 of *Listing 3.1*): only when the buffer is full (or when the amount of data to be written does not fit the buffer) file writing actually occurs. The choice of using a buffer was made as an attempt to improve performance, as it allows the aggregation of small write operations into a block size that is more efficient for the disk subsystem.

A problem which was encountered when trying to conjugate file management and multithreading is that there exists a known isolation issue affecting Pin on Windows. Specifically, it is possible for a deadlock to take place if a file is opened in a callback in the context of multithreaded applications. As a result of this issue, it is not possible to open the tracer's log file in the `OnThreadStart` callback. Pin's manual proposes to circumvent the problem by opening the file in the `main` and tagging the data with the thread ID [15]. However, this conflicts with the idea of having one file for each thread.

In order to bypass this limitation of the Pin's framework, the following strategy was employed:

1. When initializing the thread local data in `OnThreadStart`, `file_write` is set to point to a function named `file_open`.
2. The first time a thread attempts to write data to file `file_open` is invoked.

3. `file_open` carries out the following actions:
 - (a) Opens the tracer's log file (this is safe since the file is not opened in a callback)
 - (b) Sets the obtained file pointer in the thread local data (line 8 of *Listing 3.1*)
 - (c) Adds the data to be written in the buffer (which is eventually written to file if the buffer is too small to hold it)
 - (d) Sets `file_write` to point to `buf_write`, a function which is in charge of just writing data to the buffer and to file.
4. As a result of this, when the thread attempts to write to file again, `buf_write` is invoked, thus allowing the thread to just write to file without going through opening the file again.

3.2 Native APIs Tracing

Windows Native APIs are employed to call operating system services in kernel mode in a controlled way. In fact, all core Windows components, which possess direct access to hardware and services in charge of handling the computer's resources (e.g. memory), operate in kernel mode. This means that, every time a user mode application desires to carry out certain actions, like for instance starting a thread or allocating virtual memory, they must rely on kernel mode services. The Windows Native API corresponds to the system call interface of standard monolithic operating systems, such as the majority of UNIX-like systems, with the difference that in the latter case the system call interface is documented and can be utilized directly by applications. Instead, due to Windows' architecture, Windows Native APIs are concealed to the programmer by the higher level Windows (Win32) APIs [16]. User mode Windows Native APIs, which are identified by their `Nt` prefix and are exported by `ntdll`, have caught the attention of malware writers since they are seen as a way of bypassing the documented APIs with the objective of performing a series of actions without being discovered [4]. For this reason it is a good idea to also trace them, in addition to the ordinary Windows APIs.

Pin provides a set of APIs aimed at assisting the extraction of information relative to the system calls made by the pinned application, also including information relative to Windows Native APIs. In particular, in Blue Pill (and consequently in BlueTracer) `PIN_AddSyscallEntryFunction` and `PIN_AddSyscallExitFunction` were used for this purpose. In fact, these allow to register notification functions which are called immediately before and after the execution of a system call [15]. In BlueTracer, the function in charge of gathering Native API information before execution is `TraceSysCallEntry` while the one responsible for collecting Native API information after execution is `TraceSysCallExit`. They have been both embedded in Blue Pill's notification functions and their overall structure is detailed in *Figure 3.2*.

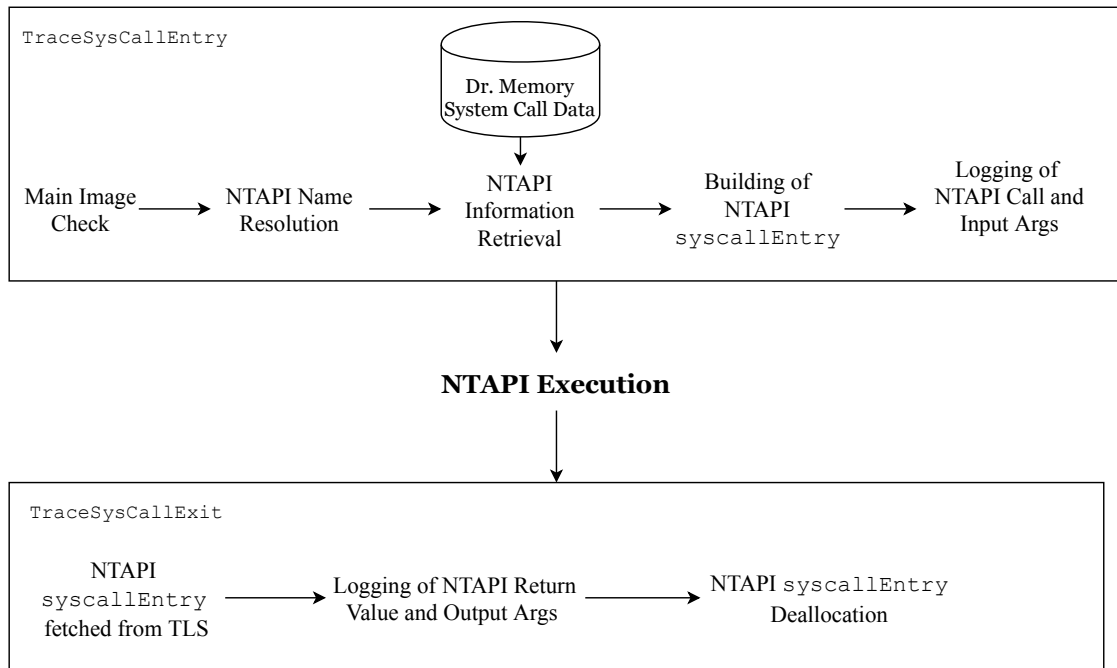


Figure 3.2. *Native APIs Tracing Workflow*

Following *Figure 3.2*, let us now thoroughly analyze the main steps which take place when Native APIs are traced.

3.2.1 Main Image Check

In order to filter the logged information relative to Native APIs, BlueTracer allows the analyst to decide, through the use of a boolean parameter (**MainImage**) in the configuration file, whether or not only Native APIs called invoked directly from the main executable of the pinned application should be traced. Therefore, as it can be seen in *Figure 3.2*, the first thing which is done in **TraceSysCallEntry**, assuming that the **MainImage** configuration parameter has been set to **true**, is to determine if the Native API call is taking place directly from the main executable.

To this end, Pin's **IMG** APIs are employed, where in Pin an **IMG** represents all the data structures relative to binaries and shared libraries [15]. Specifically, in Blue Pill, **IMG_AddInstrumentFunction** is utilized to register a callback which is invoked each time an image is loaded. Inside such callback, which takes as one of the input parameters the **IMG** object representing the image being loaded, the following steps are taken:

1. **IMG_IsMainExecutable** is employed to determine if the image being loaded is the main executable of the pinned application.
2. If this is the case, **IMG_HighAddress** and **IMG_LowAddress** are invoked. These two APIs return the highest address and the lowest address respectively of any code or data loaded by the image corresponding the the **IMG** object they take as input. By employing these APIs is therefore possible to determine the address range relative to the main executable.

Having obtained the main executable address range in this way, it is then employed in **TraceSysCallEntry** to learn if the Native API call is occurring directly from the main image. Internally, before entering kernel mode every Native API executes some common code, i.e. each Native API stores its ordinal in the **eax** register and invokes **KiFastSystemCall**, where the **sysenter** instruction is used to actually enter kernel mode. When **sysenter** is executed, the kernel obtains the ordinal number from **eax** and utilizes it to call the corresponding function, prior to going back to user mode [11].

In Pin, the Native API is intercepted right before **sysenter** is executed.

This could be inferred by the fact that, in `TraceSysCallEntry`, the instruction pointer (obtainable via the `PIN_GetContextReg` API) contains `sysenter`'s address. In light of this, in addition to the fact that each Native API calls `KiFastSystemCall` without setting a stack frame, the application's stack during the execution of `TraceSysCallEntry` is in the following state (*Figure 3.3*).

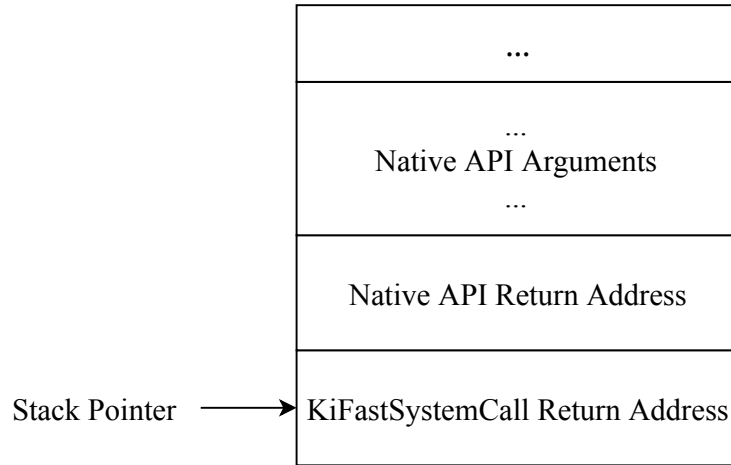


Figure 3.3. *Stack before `sysenter` execution*

Having outlined the general situation, it is finally possible to describe how the main image check is carried out:

1. The first check involves discovering whether or not `sysenter` is called directly from the main executable. As Pin intercepts Native APIs right before the execution of `sysenter`, this can be trivially done by checking if the instruction pointer value obtained in `TraceSysCallEntry` falls within the main executable memory address range.
2. Secondly, it is required to examine if the Native API return address belongs to the main executable. As a result of what shown in *Figure 3.3* such return address is obtained by adding 4 to the stack pointer and retrieving the pointed value, where the stack pointer can be obtained in Pin through the use of `PIN_GetContextReg`. The resulting value is then, once more, compared with the main executable memory address range.

3.2.2 Native API Name Resolution

Previously, it was stated how it is possible to register a callback function to be executed before Native APIs through the use of `PIN_AddSyscallEntryFunction`. Such callback functions, named `SYSCALL_ENTRY_CALLBACKs`, receive a set of parameters, including `ctx`, the application's register state immediately before the system call execution, and `std`, the system call standard. In Blue Pill, these two parameters are used to invoke `PIN_GetSyscallNumber` inside its `SYSCALL_ENTRY_CALLBACK`. Such API returns the number (ID) of the Native API to be executed in the provided context [15].

Blue Pill employs a mechanism which allows it to obtain the Native API name from its number, since the hooking functions it employs are indexed by name and not by ordinal. This choice was made because the identifiers vary depending on the Windows version, even among different Service Pack versions. For the same reason, the information needed by BlueTracer to correctly log Native APIs is also indexed by Native API name and, consequently, BlueTracer also utilizes the aforementioned Native API name resolution mechanism.

The idea is to create an array, named `syscallIDs`, where Native API names are indexed by their IDs, by parsing `ntdll`'s export information. The export data of a PE is stored in the `IMAGE_EXPORT_DIRECTORY` structure located in the header. In our case, the most relevant fields of this structure are:

- The `AddressOfFunctions` array, which contains RVAs¹ pointing to the actual exported functions and is indexed by an export ordinal.
- The `AddressOfNames` array, which is an array of 32-bit RVAs pointing to symbol strings.
- The `AddressOfNameOrdinals` array, which is an array of 16-bit ordinals existing in parallel with `AddressOfNames`, i.e. they possess the same number of elements and there is a direct relation between equivalent indices [17].

¹A RVA (Relative Virtual Address) is essentially an offset within the PE image in memory

With this in mind, `syscallIDs` is built by carrying out the following actions for every element `iName` in `AddressOfNames`:

1. The corresponding `AddressOfNameOrdinals` element is retrieved, i.e. the one with the same index. Let us call this value `iOrdinal`.
2. `iOrdinal` is used to index in `AddressOfFunctions`. This time the obtained value is the RVA pointing to the exported function.
3. The pointed function is expected to begin `mov eax, syscall_number`. As a result of this, the first byte should be `B8h` and the syscall number can be obtained by considering the next four bytes.

By following the above procedure, every element in `AddressOfNames` can be therefore mapped to the corresponding identifier.

3.2.3 Native API Information Representation and Retrieval

When tracing Native APIs, it is wanted to record as much data as possible as well as correctly formatting the arguments' values. To do this, it is required to have access to some kind of source of Native API related information, which can assist the tracer in the logging activity by, for example, providing the number of arguments a Native API takes as input, listing the arguments' types and differentiating between input and output arguments. In BlueTracer, this information was adapted from the Native API data provided by Dr.Memory [7], a memory monitoring tool based on the DynamoRIO DBI framework.

For each Native API, the information related to it is contained in a `struct` of type `syscall_info_t` (*Listing 3.2*). Most fields of such `struct` are self-explanatory. In particular, `num` is a `struct` storing two values indicating the system call number; these are filled in dynamically by Dr.Memory but are not needed by BlueTracer. Furthermore, the `flags` field is utilized to notify whether or not all the details of the Native API are known, as most Native APIs are undocumented. Undoubtedly though, a major role is played by the `arg` array, which is made up by `sysinfo_arg_t` structs (*Listing 3.2*) containing the data of the Native API arguments. This array

is initialized with size `MAX_ARGS_IN_ENTRY` (i.e. 18), since a Native API can have at most 18 arguments.

```

1  typedef struct _syscall_info_t {
2
3      drsys_sysnum_t num;           // Native API ID
4      const char *name;           // Native API Name
5      uint flags;                 // SYSINFO_ flags
6      uint return_type;           // Return type
7      int arg_count;              // Number of arguments
8
9      // Array of arguments
10     sysinfo_arg_t arg[MAX_ARGS_IN_ENTRY];
11
12     ...
13 } syscall_info_t;

```

Listing 3.2. struct containing Native API related information

```

1  typedef struct _sysinfo_arg_t {
2      int param;                  // Parameter Ordinal
3      int size;                   // Size
4      uint flags;                 // SYSARG_ flags
5      int type;                   // Type
6      const char *type_name;      // Symbolic Name of the arg Type
7  } sysinfo_arg_t;

```

Listing 3.3. struct containing information associated to a Native API argument

In `sysinfo_arg_t`, other than `param` and `size`, that are straightforward, the following fields are also present:

- `flags`, which stores an OR of flags describing the argument's characteristics.

The most important ones are:

- `SYSARG_READ (R)` : input argument.
- `SYSARG_WRITE (W)` : output argument.

- `SYSARG_INLINED` : non-memory argument, i.e the whole value is in parameter slot.
 - `SYSARG_HAS_TYPE` (HT) : argument with a type specifier. In fact, a non `SYSARG_INLINED` argument is by default of type `struct` (`DRSYS_TYPE_STRUCT`), unless specified otherwise by `SYSARG_HAS_TYPE`.
- `type`, which is an `enum` value indicating the data type of the parameter
 - `type_name`, a string indicating the symbolic name of the `arg` type, which is typically filled dynamically based on the `type` value. In case the argument value is a named constant, the `type_name` field contains the name of the enumeration the constant belongs to. In fact, for each one of these enumerations, BlueTracer has access to a `struct` specific for that enumeration containing the constant values and their corresponding name. This information was also provided by Dr.Memory and is employed to translate constants to the appropriate name, as it will be explained in the next section.

To make things clearer, an instance of `syscall_info_t` for the Native API `NtAllocateVirtualMemory` was provided in the listing below (*Listing 3.4*).

```

1      { { 0,0 }, "NtAllocateVirtualMemory", OK, RNTST, 6,
2      {
3          { 0, WIN_SIZE(HANDLE), SYSARG_INLINED, DRSYS_TYPE_HANDLE },
4          { 1, WIN_SIZE(PVOID), R | WR | HT, DRSYS_TYPE_POINTER },
5          { 2, WIN_SIZE(ULONG), SYSARG_INLINED, DRSYS_TYPE_UNSIGNED_INT },
6          { 3, WIN_SIZE(ULONG), R | WR | HT, DRSYS_TYPE_UNSIGNED_INT },
7          { 4, WIN_SIZE(ULONG), SYSARG_INLINED, DRSYS_TYPE_UNSIGNED_INT, "MEM_COMMIT" },
8          { 5, WIN_SIZE(ULONG), SYSARG_INLINED, DRSYS_TYPE_UNSIGNED_INT, "PAGE_NOACCESS" },
9      }

```

Listing 3.4. Instance of `syscall_info_t` relative to `NtAllocateVirtualMemory`

In 3.2.2 it was explained how BlueTracer obtains a Native API name from its ID. After that, as shown in *Figure 3.2*, there must be a way to, given a specific Native API name, retrieve the corresponding `syscall_info_t` struct. This is done by building an hash map during the tool's initialization which maps the Native API

name to the appropriate `struct`. Such idea is also adopted for named constants enumerations since a hash map is also constructed with the objective of associating the name of a named constant enumeration to the corresponding `struct`.

3.2.4 Native API Logging

With reference to *Figure 3.2*, let us consider the point in `TraceSysCallEntry` where the `syscall_info_t struct` relative to the Native API being traced was successfully retrieved. Before describing how the logging of "known" Native APIs is performed, it is worth mentioning that, in case a `struct` describing a certain Native API is not found, BlueTracer just logs the name of the Native API and the input value of a user-specified number of arguments (4 by default).

The first step towards logging "known" Native APIs is allocating and initializing a `syscall_tracer struct` (*Listing 3.5*) representing the specific Native API being traced. This `struct` is built starting from the data present in the previously retrieved `syscall_info_t struct` and it is at the heart of the logging of the Native API.

```

1  typedef struct _syscall_tracer {
2      ADDRINT syscall_number;          // Native API ID
3      const char * syscall_name;      // Native API Name
4      int argcount;                   // Number of Arguments
5
6      // Array of arguments
7      drsys_arg_t arguments[MAX_ARGS_IN_ENTRY];
8
9      drsys_arg_t retval;              // Return value
10     int syscall_counter;              // Native API Counter
11 } syscall_tracer;

```

Listing 3.5. `struct` representing the Native API being traced

Even for `syscall_tracer`, each field is self-explanatory. The only one which was not encountered before is the `syscall_counter` field. The rationale behind this field is that, in the log, each traced call has a unique integer associated

to it, utilized to group together the information relative to the call in post-processing. The next identifier to be assigned is stored in the `call_number` field of the thread local storage (*Listing 3.1*). Therefore, `syscall_counter` simply contains the unique identifier associated to the Native API in the log. The arguments of a specific Native API are represented via the `drsys_arg_t` struct (*Listing 3.6*).

```

1  typedef struct _drsys_arg_t {
2
3      // Whether operating pre-call (if true) or post-system call (if false)
4      bool pre;
5
6      int ordinal;           // Parameter Ordinal
7      int size;             // Size
8      uint flags;           // SYSARG_ flags
9      int type;             // Type
10     const char *type_name; // Symbolic Name of the arg Type
11
12     // String describing the symbolic name of a named constant's enum
13     const char *enum_name;
14
15     // Argument value
16     ADDRINT value;
17
18 } drsys_arg_t

```

Listing 3.6. struct representing the an argument of the Native API being traced

`drsys_arg_t`'s fields are mostly the same of `sysinfo_arg_t`, with three additions:

- `pre`, which is a boolean flag set to `true` when the logging is occurring before the Native API is executed (i.e. in `TraceSysCallEntry` and set to `false` when the logging is occurring after the Native API is executed (i.e. in `TraceSysCallExit`).
- `enum_name`, a string containing the symbolic name of the enumeration a named constant belongs to. In most cases, this is the symbolic name associated to

the first constant of the enumeration. Such field was introduced to decouple type name and enumeration name.

- **value**, an `ADDRINT`² containing the value of the argument. This is obtained through the `PIN_GetSyscallArgument` API, which takes as input the context before the execution of the system call, the system call standard and the ordinal number of the argument whose value is requested.

As previously mentioned in section 3.1.1, each thread possesses its own log file, whose pointer resides in the thread local storage, and writes on it by means of the `file_write` function, also located in the thread local storage (*Listing 3.1*). The first Native API component being logged is its name and this is done with the following format:

```
~~[System ID of thread]~~ [Call Counter] [Image Name]![Native API Name]
```

In the context of Native APIs, the name of the image is always `ntdll`, but this is not the case when tracing APIs.

Argument logging is done via the `print_arg` function (*Listing 3.7*).

²The `ADDRINT` type is defined by Pin and represents a memory address

```
1 static void print_arg(drsys_arg_t* curr_arg, bluepill_tls* tdata, uint syscall_counter) {
2
3     ...
4
5     // Constant Resolution
6     if (curr_arg->enum_name != NULL) {
7         if (get_arg_symname(curr_arg, tdata))
8             return;
9     }
10
11     switch (curr_arg->type) {
12         case DRSYS_TYPE_VOID:      print_simple_value(curr_arg, true, tdata); break;
13         case DRSYS_TYPE_POINTER:   print_simple_value(curr_arg, true, tdata); break;
14         case DRSYS_TYPE_BOOL:      print_simple_value(curr_arg, false, tdata); break;
15         case DRSYS_TYPE_INT:       print_simple_value(curr_arg, false, tdata); break;
16         case DRSYS_TYPE_SIGNED_INT: print_simple_value(curr_arg, false, tdata); break;
17         case DRSYS_TYPE_UNSIGNED_INT: print_simple_value(curr_arg, false, tdata); break;
18         case DRSYS_TYPE_HANDLE:    print_simple_value(curr_arg, false, tdata); break;
19         case DRSYS_TYPE_NTSTATUS:  print_simple_value(curr_arg, false, tdata); break;
20         case DRSYS_TYPE_ATOM:      print_simple_value(curr_arg, false, tdata); break;
21     default: {
22         if (curr_arg->value == 0) {
23             (tdata->file_write)(tdata->threadid, tdata->buffer, tdata->OutFile, "<null>");
24         }
25         else if (!TEST(SYSARG_READ, curr_arg->flags)) {
26             (tdata->file_write)(tdata->threadid, tdata->buffer,
27                 tdata->OutFile, PFX, curr_arg->value);
28         }
29
30         else {
31             if (!print_known_compound_type(curr_arg, tdata))
32                 (tdata->file_write)(tdata->threadid, tdata->buffer, tdata->OutFile, "<NYI>");
33         }
34
35     }
36 }
```

Listing 3.7. print_arg function

3.3 API

In tracing APIs most of the idea seen in Native API tracing also apply, with a few differences. -> Same diagram as NTAPI with 2 sources AND shadow stack

3.3.1 Shadow Stack

3.3.2 Performance

3.4 Callback and APC

Chapter 4

Experimental Results

Chapter 5

Conclusions and Future Developments

Bibliography

- [1] Ulrich Bayer and Christopher Krügel. “TTAnalyze : A Tool for Analyzing Malware”. In: 2005.
- [2] J. Berdajs and Z. Bosnić. “Extending Applications Using an Advanced Approach to DLL Injection and API Hooking”. In: *Softw. Pract. Exper.* 40.7 (June 2010), pp. 567–584. ISSN: 0038-0644. DOI: 10.1002/spe.v40:7. URL: <http://dx.doi.org/10.1002/spe.v40:7>.
- [3] Armin Buescher, Felix Leder, and Thomas Siebert. “Banksafe Information Stealer Detection Inside the Web Browser”. In: *Proceedings of the 14th International Conference on Recent Advances in Intrusion Detection*. RAID’11. Menlo Park, CA: Springer-Verlag, 2011, pp. 262–280. ISBN: 978-3-642-23643-3. DOI: 10.1007/978-3-642-23644-0_14. URL: http://dx.doi.org/10.1007/978-3-642-23644-0_14.
- [4] Geoff Chappell. *Native API Functions*. 2017. URL: <https://www.geoffchappell.com/studies/windows/win32/ntdll/api/native.htm>.
- [5] *Cisco 2018 Annual Cybersecurity Report*. 2018. URL: https://www.cisco.com/c/dam/m/hu_hu/campaigns/security-hub/pdf/acr-2018.pdf.
- [6] Daniele Cono D’Elia, Emilio Coppa, and Camil Demetrescu. *The DBI Blue Pill: Practical Analysis of Evasive Malware*.
- [7] *Dr.Memory*. URL: <http://drmemory.org/>.
- [8] *Dynamic Binary Instrumentation*. 2007. URL: <http://uninformed.org/index.cgi?v=7&a=1&p=3>.
- [9] *Dynamic Instrumentation Tool Platform*. URL: <http://www.dynamorio.org/>.

- [10] Manuel Egele et al. “A Survey on Automated Dynamic Malware-analysis Techniques and Tools”. In: *ACM Comput. Surv.* 44.2 (Mar. 2008), 6:1–6:42. ISSN: 0360-0300. DOI: 10.1145/2089125.2089126.
- [11] *Intercepting all System Calls by Hooking KiFastSystemCall*. 2015. URL: <https://www.malwaretech.com/2015/04/intercepting-all-system-calls-by.html>.
- [12] *Internet Security Threat Report, vol. 23*. 2018. URL: <https://www.symantec.com/content/dam/symantec/docs/reports/istr-23-2018-en.pdf>.
- [13] Ivo Ivanov. *API hooking revealed*. 2002. URL: <https://www.codeproject.com/Articles/2082/API-hooking-revealed>.
- [14] Syed Zainudeen Mohd Shaid and Mohd Maarof. “In memory detection of Windows API call hooking technique”. In: (Aug. 2015), pp. 294–298.
- [15] *Pin - A Dynamic Binary Instrumentation Tool*. 2012. URL: <https://software.intel.com/en-us/articles/pin-a-dynamic-binary-instrumentation-tool>.
- [16] Mark Russinovich. *Inside the Native API*. 1998. URL: <http://persephone.cps.unizar.es/~spd/pub/windows/ntdll.htm>.
- [17] Michael Sikorski and Andrew Honig. *Practical Malware Analysis: The Hands-On Guide to Dissecting Malicious Software*. 1st. San Francisco, CA, USA: No Starch Press, 2012. ISBN: 1593272901, 9781593272906.
- [18] Sherri Sparks, Shawn Embleton, and Cliff C. Zou. “WINDOWS ROOTKITS A GAME OF HIDE AND SEEK”. In: *Handbook of Security and Networks*. 2011, pp. 345–368. DOI: 10.1142/9789814273046_0014. eprint: https://www.worldscientific.com/doi/pdf/10.1142/9789814273046_0014. URL: https://www.worldscientific.com/doi/abs/10.1142/9789814273046_0014.
- [19] Dafydd Stuttard et al. *Attack and Defend Computer Security Set*. 1st. Wiley Publishing, 2014. ISBN: 111890673X, 9781118906736.